

## Time in Distributed Systems

In distributed systems, multiple independent computers (or nodes) work together to achieve a common goal. Since these systems often span across different physical locations, managing time effectively becomes crucial for coordination and consistency. Below are key concepts regarding time in distributed systems:

### 1. Challenges of Time in Distributed Systems

- **No Global Clock:** Unlike a single machine, distributed systems don't have a shared global clock. Each node has its own local clock, and these clocks can drift apart.
- **Clock Synchronization:** Ensuring that the clocks across different machines are synchronized is a fundamental challenge, as desynchronized clocks can lead to issues like incorrect ordering of events or inconsistency in data.
- **Latency and Communication Delays:** Network delays between nodes can cause the perception of time differences, adding complexity to time-sensitive operations.

### 2. Clock Synchronization Algorithms

To overcome the challenges of having different local clocks, distributed systems use synchronization algorithms to keep clocks relatively in sync:

- **NTP (Network Time Protocol):** A common protocol used to synchronize time across computers over a network. It uses a hierarchical system of time sources, with each node communicating with a trusted server (stratum 0) to adjust its local clock.
- **Berkeley Algorithm:** A decentralized algorithm used for synchronizing clocks in a distributed system. It works by selecting one node as the coordinator, which requests the time from all other nodes, calculates the average time, and sends the correction to each node.
- **Cristian's Algorithm:** A client-server-based method in which a client requests the current time from a time server and adjusts its own clock based on the server's timestamp, considering the communication delay.

### 3. Logical Time vs. Physical Time

- **Physical Time:** Refers to the actual time represented by the system's clock, typically synchronized with real-world time. However, as mentioned earlier, physical time can vary from one node to another due to clock drift.
- **Logical Time:** A logical clock is a mechanism that allows the ordering of events without relying on physical time. Logical time is used to preserve the causal relationship between events in the system.
  - **Lamport Timestamps:** One of the most commonly used logical clocks, which assigns a number (timestamp) to each event and increments it with each new event. It helps in determining the order of events but does not provide real-time synchronization.

- **Vector Clocks:** An extension of Lamport timestamps, vector clocks keep track of causal relationships by maintaining a vector of logical clocks at each node. This allows detecting causality between events more accurately, such as whether one event happened before, after, or concurrently with another.

#### 4. Time in Consensus Protocols

In distributed systems, especially when coordinating across nodes to achieve consistency (such as in distributed databases or blockchain systems), time plays a key role in achieving consensus.

- **Paxos:** A consensus algorithm used to ensure agreement across nodes on a single value. While Paxos doesn't directly rely on real-world time, having synchronized clocks helps avoid situations where nodes are working off different timelines and lead to inconsistencies.
- **Raft:** Another consensus algorithm that is easier to understand than Paxos and relies on leader election and log replication. Time synchronization can help nodes know the current leader and avoid conflicting operations.

#### 5. Event Ordering in Distributed Systems

In the absence of synchronized physical clocks, determining the order of events is critical. Various models can help:

- **Happens-Before Relation:** In a distributed system, an event A is said to "happen before" event B if the occurrence of A influences the occurrence of B. This is fundamental for establishing causality, ensuring events occur in the right order.
- **Vector Clocks (Continued):** As mentioned earlier, vector clocks are particularly useful for determining if events are causally related, concurrent, or if one happened before the other. This helps in maintaining consistency in the system.

#### 6. Clock Drift and Skew

- **Clock Drift:** Refers to the gradual difference between two clocks over time. Even with synchronization algorithms, clocks can drift slightly due to factors like temperature and hardware differences. This drift can cause issues in time-sensitive applications.
- **Clock Skew:** Refers to the difference between two clocks at a given point in time. Synchronization algorithms try to minimize skew, but it can still occur in large-scale distributed systems.

#### 7. Use of Time in Fault Tolerance

- **Event Logging and Recovery:** Time is critical when recording events in logs or recovering from failures. For instance, in the event of a node failure, having accurate timestamps (either physical or logical) can help in determining the point of failure and recovering to a consistent state.
- **Timeouts:** In protocols like distributed locking or leader election, timeouts are essential to decide when a node should take action or wait for a response.

#### 8. Clock Synchronization in Large-Scale Systems

- **GPS Time:** In large-scale systems spanning multiple geographical locations, GPS time is often used as a standard reference to synchronize nodes. GPS satellites provide highly accurate time signals that can be utilized for precise synchronization.

# How clock Synchronization Algorithms work in distributed systems

## 1. Network Time Protocol (NTP)

### How it works:

- NTP synchronizes a client's clock with a time server by calculating the round-trip time (RTT) and adjusting the client's clock.
- The basic idea is to account for the time delay between the client and server.

### Formula:

1. **T1**: The time when the request is sent from the client to the server.
2. **T2**: The time when the request arrives at the server.
3. **T3**: The time when the response is sent from the server to the client.
4. **T4**: The time when the response is received by the client.

The client's clock adjustment is calculated as:

$$\text{Offset} = \frac{(T2 - T1) + (T3 - T4)}{2}$$

Where:

- **Offset** is the time difference between the server's clock and the client's clock.
- **Delay** is the total round-trip delay:

$$\text{Delay} = \frac{(T3 - T2) + (T4 - T1)}{2}$$

### Example:

- Suppose the following times were recorded:
  - T1=10:00:00 (client sends request)
  - T2=10:00:05 (server receives request)
  - T3=10:00:10 (server sends response)
  - T4=10:00:15 (client receives response)

The **Offset** would be:

$$\text{Offset} = \frac{(10:00:05 - 10:00:00) + (10:00:10 - 10:00:15)}{2} = \frac{5 + (-5)}{2} = 0$$

The **client's clock** would need no adjustment in this case, but if there was a non-zero offset, the client would adjust accordingly.

## 2. Berkeley Algorithm

### How it works:

- In the Berkeley algorithm, a coordinator node asks all other nodes for their current time, calculates the average, and sends the time adjustment back to each node to synchronize their clocks.
- The coordinator determines the time offset for each node.

### Formula:

1. Let the **local time** of node  $i$  be  $T_i$ .
2. The **coordinator's time**  $TC$  is taken as a reference, and the **average time** of all nodes is calculated:

$$\text{Average Time} = \frac{\sum_{i=1}^n T_i}{n}$$

3. Each node adjusts its clock by the difference between its time and the average time:

Adjustment for node  $i = \text{Average Time} - T_i$

### Example:

Assume there are 3 nodes, and their current times are:

- Node 1:  $T_1 = 10:00:00$
- Node 2:  $T_2 = 10:00:10$
- Node 3:  $T_3 = 10:00:20$

The average time is:

$$\text{Average Time} = \frac{10:00:00 + 10:00:10 + 10:00:20}{3} = 10:00:10$$

Each node would adjust its clock:

- Node 1 adjusts by  $10:00:10 - 10:00:00 = 10$  seconds ahead.
- Node 2 adjusts by  $10:00:10 - 10:00:10 = 0$  seconds (no adjustment).
- Node 3 adjusts by  $10:00:10 - 10:00:20 = -10$  seconds (10 seconds behind).

### 3. Cristian's Algorithm

#### How it works:

- Cristian's algorithm is a simpler, client-server-based synchronization protocol where the client sends a request to a time server, and the server responds with its current time.
- The client estimates the round-trip delay and adjusts its clock by half of that delay.

#### Formula:

1. **T1**: Time when the client sends the request.
2. **T2**: Time when the server receives the request.
3. **T3**: Time when the server sends the response back to the client.
4. **T4**: Time when the client receives the response.

The round-trip delay is:

$$\text{Round-trip Delay} = (T4 - T1) - (T3 - T2)$$

The clock offset is:

$$\text{Offset} = \frac{(T2 - T1) + (T3 - T4)}{2}$$

#### Example:

Assume the following times:

- T1=10:00:00 (client sends request)
- T2=10:00:02 (server receives request)
- T3=10:00:08 (server sends response)
- T4=10:00:10 (client receives response)

The **Offset** is:

$$\text{Offset} = \frac{(10:00:02 - 10:00:00) + (10:00:08 - 10:00:10)}{2} = \frac{2 - 2}{2} = 0$$

The **client's clock** does not need to be adjusted here.

## 5. Vector Clocks

### How it works:

- Vector clocks extend Lamport's clocks by allowing each process to maintain a vector of timestamps, one for each process in the system.
- When an event occurs, each node increments its entry in the vector. When a message is sent, it includes the entire vector, allowing the recipient to determine causality.

### Formula:

Let  $V_i(t)$  be the vector clock of process  $i$  at time  $t$ .

1. When an event occurs at node  $i$ , increment its entry in the vector:

$$V_i(t) = (V_1(t), V_2(t), \dots, V_n(t)) \quad \text{where } V_i(t) = V_i(t-1) + 1$$

2. When a message is sent from node  $i$  to node  $j$ , it sends  $V_i(t)$  along with the message.
3. When node  $j$  receives the message, it updates its vector clock:

$$V_j(t) = \max(V_j(t-1), V_i(t)) \quad \text{and then increment its own vector entry: } V_j(t) = V_j(t) + 1$$

### Example:

- Node A's vector clock is  $V_A = (3, 2, 4)$
- Node B's vector clock is  $V_B = (2, 3, 4)$
- If Node A sends a message to Node B, Node B updates its clock as:

$$V_B = \max(V_B, V_A) + 1 = \max((2, 3, 4), (3, 2, 4)) + 1 = (3, 3, 4) + 1 = (3, 3, 5)$$